

## Vector clock

A *vector clock* is a tuple of  $N$  logical clocks, one per thread:

$$(c_0, c_1, \dots, c_{N-1})$$

Equivalently: vector clock is a mapping from *Threads* to natural numbers

If  $V$  is a vector clock,  $V(t)$  denotes logical clock of thread  $t$

Set of all vector clocks:  $VC$

## Vector clock

**Bottom vector clock:**  $\perp = (0, 0, \dots, 0)$

Every thread has logical clock 0

**Partial order on vector clocks:**  $V_1 \sqsubseteq V_2$  if and only if  $\forall t. V_1(t) \leq V_2(t)$

One vector clock is “less-than or equal to” another vector clock if logical clocks are pointwise less-than or equal

**Join of vector clocks:**  $V_1 \sqcup V_2 =$  pointwise maximum of  $V_1, V_2$

**Increment:**  $inc_t(V) = V[t \mapsto V(t) + 1]$

Same as  $V$ , but logical clock for  $t$  is incremented

## Purpose of per-thread vector clocks

$C : \text{Threads} \rightarrow VC$	Each <i>thread</i> has a vector clock
-------------------------------------	---------------------------------------

$C_t$  - shorthand for  $C(t)$

$C_t$  represents what thread  $t$  knows about the logical clocks of all threads

If I am thread  $t$ , then:

- $C_t(t)$  is *my* logical clock. It will always be **positive**
- For  $u \neq t$ ,  $C_t(u) = z$  means “I know that thread  $u$ ’s logical clock is at least  $z$ ”

When  $t$  **acquires a lock**,  $t$  gets information about the logical clocks of threads who previously held the lock

## Purpose of per-lock vector clocks

$L : \text{Locks} \rightarrow VC$	Each lock has a vector clock
-----------------------------------	------------------------------

$L_m$  - shorthand for  $L(m)$

$L_m$  represents the logical clock each thread had last time it **released**  $m$

For a lock  $m$  and thread  $t$ :

Remember that a thread's true logical clock value will always be positive

- $L_m(t) = 0$  means that  $t$  has never released  $m$
- Otherwise  $L_m(t)$  was  $t$ 's logical clock last time  $t$  released  $m$

## Purpose of per-location read vector clocks

$R : \text{Locations} \rightarrow VC$ Each <i>location</i> has a <i>read</i> vector clock
---

$R_x$  - shorthand for  $R(x)$

$R_x$  represents the logical clock each thread had last time it **read from**  $x$

For a location  $x$  and thread  $t$ :

- $R_x(t) = 0$  means that  $t$  has never read from  $x$
- Otherwise  $R_x(t)$  was  $t$ 's logical clock last time  $t$  read from  $x$

## Purpose of per-location write vector clocks

$W : \text{Locations} \rightarrow VC$ Each <i>location</i> has a <i>write</i> vector clock
--

$W_x$  - shorthand for  $W(x)$

$W_x$  represents the logical clock each thread had last time it **wrote to**  $x$

For a location  $x$  and thread  $t$ :

- $W_x(t) = 0$  means that  $t$  has never written to  $x$
- Otherwise  $W_x(t)$  was  $t$ 's logical clock last time  $t$  wrote to  $x$

## Initial analysis state

- $C = (inc_0(\perp), inc_1(\perp), \dots, inc_{N-1}(\perp))$

Thread  $t$  knows its logical clock is 1, and thinks all other threads' logical clocks are 0

- $L = \lambda m . \perp$

No thread has unlocked any mutex

- $R = \lambda x . \perp$
- $W = \lambda x . \perp$

No thread has read from or written to any location

## Vector clock algorithm state

Recall:

- $N$  threads, starting from 0
- Locks: available locks for synchronisation
- *Locations*: possibly-shared locations, on which races could occur

Algorithm state:  $(C, L, R, W)$ :

- $C : \text{Threads} \rightarrow VC$       Each *thread* has a vector clock
- $L : \text{Locks} \rightarrow VC$       Each *lock* has a vector clock
- $R : \text{Locations} \rightarrow VC$       Each *location* has a *read* vector clock
- $W : \text{Locations} \rightarrow VC$       Each *location* has a *write* vector clock